

# MIS780 Advanced AI For Business - Assignment 2 - T2 2022

## Demonstrative Example Number 3: Recurrent Neural network - MasterCard Stock Price Prediction Using LSTM & GRU

**Student Name:** SHIFAT ADBULLAH

**Student ID:** 220144552

### Table of Content

1. Data Description
2. Data Preprocessing
3. Model Construction
4. Model Execution
5. Experiments Report

### Executive Summary

RNN remembers past inputs due to an internal memory which is useful for predicting stock prices, generating text, transcriptions, and machine translation. In the traditional neural network, the inputs and the outputs are independent of each other, whereas the output in RNN is dependent on prior elements within the sequence. Recurrent networks also share parameters across each layer of the network.

The business problem which we are trying to address here is Predicting the MasterCard Stock Price Using LSTM & GRU.

First we will analyze data, preprocess the data to train it on advanced RNN models, and finally evaluate the results.

### 1. Data Description

First we will import the MasterCard dataset by adding the Date column to the index and converting it to DateTime format. We will also drop irrelevant columns from the dataset as we are only interested in stock prices, volume, and date. The dataset has Date as index and Open, High, Low, Close, and Volume as columns.

*# Importing the libraries*

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
from sklearn.preprocessing import MinMaxScaler
```

```

from sklearn.metrics import mean_squared_error

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, Dropout, GRU,
Bidirectional
from tensorflow.keras.optimizers import SGD
from tensorflow.random import set_seed

set_seed(200)
np.random.seed(200)

dataset = pd.read_csv(
    "Mastercard_stock_history.csv", index_col="Date",
    parse_dates=["Date"]
).drop(["Dividends", "Stock Splits"], axis=1)
print(dataset.head())

```

	Open	High	Low	Close	Volume
Date					
2006-05-25	3.748967	4.283869	3.739664	4.279217	395343000
2006-05-26	4.307126	4.348058	4.103398	4.179680	103044000
2006-05-30	4.183400	4.184330	3.986184	4.093164	49898000
2006-05-31	4.125723	4.219679	4.125723	4.180608	30002000
2006-06-01	4.179678	4.474572	4.176887	4.419686	62344000

```
print(dataset.describe())
```

	Open	High	Low	Close
Volume				
count	3872.000000	3872.000000	3872.000000	3872.000000
3.872000e+03				
mean	104.896814	105.956054	103.769349	104.882714
1.232250e+07				
std	106.245511	107.303589	105.050064	106.168693
1.759665e+07				
min	3.748967	4.102467	3.739664	4.083861
6.411000e+05				
25%	22.347203	22.637997	22.034458	22.300391
3.529475e+06				
50%	70.810079	71.375896	70.224002	70.856083
5.891750e+06				
75%	147.688448	148.645373	146.822013	147.688438
1.319775e+07				
max	392.653890	400.521479	389.747812	394.685730
3.953430e+08				

We use High column to train the model. We can also choose Close or Open columns for a model feature, but High makes more sense as it provides us information of how high the values of the share went on the given day. The minimum stock price is \$4.10, and the highest is \$400.5. The mean is at 105.9 and the standard deviation 107.3, which means that stocks have high variance.

## 2. Data Preprocessing

We can see that this dataset does not have any null values.

```
dataset.isna().sum()
```


```
Open      0
High      0
Low       0
Close     0
Volume    0
dtype: int64
```

Our test dataset consists of two years, from 2020 to 2022, and the rest of the dataset is used for training.

```
tstart = 2016
tend = 2019
```

```
def train_test_plot(dataset, tstart, tend):
    dataset.loc[f"{tstart}":f"{tend}", "High"].plot(figsize=(16, 4),
    legend=True)
    dataset.loc[f"{tend+1}":, "High"].plot(figsize=(16, 4),
    legend=True)
    plt.legend([f"Train (Before {tend+1})", f"Test ({tend+1} and
    beyond)"])
    plt.title("MasterCard stock price")
    plt.show()
```

```
train_test_plot(dataset, tstart, tend)
```



```
def train_test_split(dataset, tstart, tend):
    train = dataset.loc[f"{tstart}":f"{tend}", "High"].values
    test = dataset.loc[f"{tend+1}":, "High"].values
    return train, test
training_set, test_set = train_test_split(dataset, tstart, tend)
```

We will use the MinMaxScaler function to standardize our training set, which will help us avoid the outliers or anomalies

```

sc = MinMaxScaler(feature_range=(0, 1))
training_set = training_set.reshape(-1, 1)
training_set_scaled = sc.fit_transform(training_set)

def split_sequence(sequence, n_steps):
    X, y = list(), list()
    for i in range(len(sequence)):
        end_ix = i + n_steps
        if end_ix > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_ix], sequence[end_ix]
        X.append(seq_x)
        y.append(seq_y)
    return np.array(X), np.array(y)

n_steps = 60
features = 1
# split into samples
X_train, y_train = split_sequence(training_set_scaled, n_steps)

# Reshaping X_train for model
X_train = X_train.reshape(X_train.shape[0],X_train.shape[1],features)

```

### 3. Model Construction

The model consists of a single hidden layer of LSTM and an output layer. The more units will give us better results. For this experiment, we will set LSTM units to 125, tanh as activation, and set input size. Tensorflow library is user-friendly, so we don't have to create LSTM or GRU models from scratch. We will simply use the LSTM or GRU modules to construct the model. Finally, we will compile the model with an RMSprop optimizer and mean square error as a loss function.

```

# The LSTM architecture
model_lstm = Sequential()
model_lstm.add(LSTM(units=125, activation="tanh",
input_shape=(n_steps, features)))
model_lstm.add(Dense(units=1))
# Compiling the model
model_lstm.compile(optimizer="RMSprop", loss="mse")

```

```
model_lstm.summary()
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
lstm_2 (LSTM)	(None, 125)	63500

dense\_4 (Dense)

(None, 1)

126

=====  
Total params: 63,626

Trainable params: 63,626

Non-trainable params: 0  
=====

The model will train on 50 epochs with 32 batch sizes. You can change the hyperparameters to reduce training time or improve the results. The model training was successfully completed with the best possible loss.

`model_lstm.fit(X_train, y_train, epochs=50, batch_size=32)`

Epoch 1/50

30/30 [=====] - 4s 59ms/step - loss: 0.0144

Epoch 2/50

30/30 [=====] - 2s 59ms/step - loss: 0.0030

Epoch 3/50

30/30 [=====] - 2s 59ms/step - loss: 0.0016

Epoch 4/50

30/30 [=====] - 2s 60ms/step - loss: 0.0020

Epoch 5/50

30/30 [=====] - 2s 68ms/step - loss: 0.0013

Epoch 6/50

30/30 [=====] - 2s 61ms/step - loss: 0.0010

Epoch 7/50

30/30 [=====] - 2s 59ms/step - loss: 0.0012

Epoch 8/50

30/30 [=====] - 2s 60ms/step - loss: 8.5141e-04

Epoch 9/50

30/30 [=====] - 2s 59ms/step - loss: 0.0011

Epoch 10/50

30/30 [=====] - 2s 59ms/step - loss: 7.9734e-04

Epoch 11/50

30/30 [=====] - 2s 60ms/step - loss: 8.5478e-04

Epoch 12/50

30/30 [=====] - 2s 60ms/step - loss: 7.6442e-04

Epoch 13/50

30/30 [=====] - 2s 60ms/step - loss: 7.2188e-04

Epoch 14/50

30/30 [=====] - 2s 60ms/step - loss: 6.5349e-04

Epoch 15/50

30/30 [=====] - 3s 85ms/step - loss: 6.2068e-04

Epoch 16/50  
30/30 [=====] - 2s 64ms/step - loss: 5.9433e-04

Epoch 17/50  
30/30 [=====] - 2s 60ms/step - loss: 5.4611e-04

Epoch 18/50  
30/30 [=====] - 2s 59ms/step - loss: 6.4567e-04

Epoch 19/50  
30/30 [=====] - 2s 58ms/step - loss: 6.6627e-04

Epoch 20/50  
30/30 [=====] - 2s 59ms/step - loss: 4.4368e-04

Epoch 21/50  
30/30 [=====] - 2s 58ms/step - loss: 5.0636e-04

Epoch 22/50  
30/30 [=====] - 2s 59ms/step - loss: 6.4916e-04

Epoch 23/50  
30/30 [=====] - 2s 59ms/step - loss: 5.2850e-04

Epoch 24/50  
30/30 [=====] - 2s 59ms/step - loss: 4.4711e-04

Epoch 25/50  
30/30 [=====] - 3s 93ms/step - loss: 4.1572e-04

Epoch 26/50  
30/30 [=====] - 2s 59ms/step - loss: 5.6941e-04

Epoch 27/50  
30/30 [=====] - 2s 61ms/step - loss: 3.5344e-04

Epoch 28/50  
30/30 [=====] - 2s 60ms/step - loss: 4.3698e-04

Epoch 29/50  
30/30 [=====] - 2s 59ms/step - loss: 4.6632e-04

Epoch 30/50  
30/30 [=====] - 2s 60ms/step - loss: 4.7376e-04

Epoch 31/50  
30/30 [=====] - 2s 60ms/step - loss: 3.6241e-04

Epoch 32/50  
30/30 [=====] - 2s 59ms/step - loss: 4.4521e-04

04  
Epoch 33/50  
30/30 [=====] - 2s 59ms/step - loss: 4.5075e-04  
Epoch 34/50  
30/30 [=====] - 2s 59ms/step - loss: 3.9212e-04  
Epoch 35/50  
30/30 [=====] - 2s 60ms/step - loss: 4.7223e-04  
Epoch 36/50  
30/30 [=====] - 2s 58ms/step - loss: 3.8025e-04  
Epoch 37/50  
30/30 [=====] - 2s 59ms/step - loss: 3.9966e-04  
Epoch 38/50  
30/30 [=====] - 2s 58ms/step - loss: 3.5278e-04  
Epoch 39/50  
30/30 [=====] - 2s 59ms/step - loss: 3.8088e-04  
Epoch 40/50  
30/30 [=====] - 2s 60ms/step - loss: 4.0496e-04  
Epoch 41/50  
30/30 [=====] - 2s 59ms/step - loss: 3.1012e-04  
Epoch 42/50  
30/30 [=====] - 2s 60ms/step - loss: 3.3927e-04  
Epoch 43/50  
30/30 [=====] - 2s 59ms/step - loss: 3.5521e-04  
Epoch 44/50  
30/30 [=====] - 2s 61ms/step - loss: 2.9000e-04  
Epoch 45/50  
30/30 [=====] - 2s 60ms/step - loss: 3.5306e-04  
Epoch 46/50  
30/30 [=====] - 2s 60ms/step - loss: 3.3975e-04  
Epoch 47/50  
30/30 [=====] - 2s 59ms/step - loss: 3.0760e-04  
Epoch 48/50  
30/30 [=====] - 2s 61ms/step - loss: 4.0877e-04  
Epoch 49/50

```
30/30 [=====] - 3s 91ms/step - loss: 2.5505e-04
Epoch 50/50
30/30 [=====] - 2s 59ms/step - loss: 3.6095e-04
```

```
<keras.callbacks.History at 0x7f5501179b50>
```

#### 4. Model Execution

```
dataset_total = dataset.loc[:, "High"]
inputs = dataset_total[len(dataset_total) - len(test_set) -
n_steps :].values
inputs = inputs.reshape(-1, 1)
#scaling
inputs = sc.transform(inputs)


# Split into samples
X_test, y_test = split_sequence(inputs, n_steps)
# reshape
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], features)
#prediction
predicted_stock_price = model_lstm.predict(X_test)
#inverse transform the values
predicted_stock_price = sc.inverse_transform(predicted_stock_price)

def plot_predictions(test, predicted):
    plt.plot(test, color="gray", label="Real")
    plt.plot(predicted, color="red", label="Predicted")
    plt.title("MasterCard Stock Price Prediction")
    plt.xlabel("Time")
    plt.ylabel("MasterCard Stock Price")
    plt.legend()
    plt.show()

def return_rmse(test, predicted):
    rmse = np.sqrt(mean_squared_error(test, predicted))
    print("The root mean squared error is {:.2f}.".format(rmse))

plot_predictions(test_set, predicted_stock_price)
```





```
return_rmse(test_set,predicted_stock_price)
```

The root mean squared error is 9.24.

### GRU Model

We are going to keep everything the same and just replace the LSTM layer with the GRU layer to properly compare the results. The model structure contains a single GRU layer with 125 units and an output layer.

```
model_gru = Sequential()
model_gru.add(GRU(units=125, activation="tanh", input_shape=(n_steps,
features)))
model_gru.add(Dense(units=1))
# Compiling the RNN
model_gru.compile(optimizer="RMSprop", loss="mse")
```

```
model_gru.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
gru_2 (GRU)	(None, 125)	48000
dense_5 (Dense)	(None, 1)	126

Total params: 48,126  
Trainable params: 48,126  
Non-trainable params: 0

---

```
model_gru.fit(X_train, y_train, epochs=50, batch_size=32)
```

```
Epoch 1/50
30/30 [=====] - 4s 53ms/step - loss: 0.0311
Epoch 2/50
30/30 [=====] - 2s 53ms/step - loss: 0.0021
Epoch 3/50
30/30 [=====] - 2s 53ms/step - loss: 0.0016
Epoch 4/50
30/30 [=====] - 2s 76ms/step - loss: 0.0014
Epoch 5/50
30/30 [=====] - 2s 60ms/step - loss: 9.1270e-
04
Epoch 6/50
30/30 [=====] - 2s 53ms/step - loss: 9.0105e-
04
Epoch 7/50
30/30 [=====] - 2s 53ms/step - loss: 0.0010
Epoch 8/50
30/30 [=====] - 2s 51ms/step - loss: 6.8955e-
04
Epoch 9/50
30/30 [=====] - 2s 54ms/step - loss: 8.4543e-
04
Epoch 10/50
30/30 [=====] - 2s 53ms/step - loss: 7.3059e-
04
Epoch 11/50
30/30 [=====] - 2s 52ms/step - loss: 5.6285e-
04
Epoch 12/50
30/30 [=====] - 2s 53ms/step - loss: 6.5156e-
04
Epoch 13/50
30/30 [=====] - 2s 53ms/step - loss: 6.0077e-
04
Epoch 14/50
30/30 [=====] - 2s 52ms/step - loss: 5.4861e-
04
Epoch 15/50
30/30 [=====] - 2s 53ms/step - loss: 5.8007e-
04
Epoch 16/50
30/30 [=====] - 2s 52ms/step - loss: 4.2546e-
04
Epoch 17/50
```

30/30 [=====] - 2s 53ms/step - loss: 5.2514e-04  
Epoch 18/50  
30/30 [=====] - 2s 53ms/step - loss: 5.0861e-04  
Epoch 19/50  
30/30 [=====] - 2s 53ms/step - loss: 4.2881e-04  
Epoch 20/50  
30/30 [=====] - 2s 55ms/step - loss: 5.1325e-04  
Epoch 21/50  
30/30 [=====] - 2s 53ms/step - loss: 4.6646e-04  
Epoch 22/50  
30/30 [=====] - 2s 51ms/step - loss: 4.3484e-04  
Epoch 23/50  
30/30 [=====] - 2s 53ms/step - loss: 4.2782e-04  
Epoch 24/50  
30/30 [=====] - 2s 51ms/step - loss: 4.1675e-04  
Epoch 25/50  
30/30 [=====] - 2s 52ms/step - loss: 3.6416e-04  
Epoch 26/50  
30/30 [=====] - 2s 52ms/step - loss: 3.9331e-04  
Epoch 27/50  
30/30 [=====] - 2s 54ms/step - loss: 3.2801e-04  
Epoch 28/50  
30/30 [=====] - 2s 52ms/step - loss: 4.7050e-04  
Epoch 29/50  
30/30 [=====] - 2s 53ms/step - loss: 3.1620e-04  
Epoch 30/50  
30/30 [=====] - 2s 52ms/step - loss: 3.9095e-04  
Epoch 31/50  
30/30 [=====] - 2s 52ms/step - loss: 3.9653e-04  
Epoch 32/50  
30/30 [=====] - 2s 53ms/step - loss: 3.4996e-04  
Epoch 33/50  
30/30 [=====] - 2s 53ms/step - loss: 3.4088e-04

Epoch 34/50  
30/30 [=====] - 2s 52ms/step - loss: 3.3508e-04  
Epoch 35/50  
30/30 [=====] - 2s 52ms/step - loss: 4.1599e-04  
Epoch 36/50  
30/30 [=====] - 2s 52ms/step - loss: 2.9602e-04  
Epoch 37/50  
30/30 [=====] - 2s 52ms/step - loss: 3.4013e-04  
Epoch 38/50  
30/30 [=====] - 2s 53ms/step - loss: 3.0841e-04  
Epoch 39/50  
30/30 [=====] - 2s 53ms/step - loss: 3.1379e-04  
Epoch 40/50  
30/30 [=====] - 2s 54ms/step - loss: 3.3423e-04  
Epoch 41/50  
30/30 [=====] - 2s 53ms/step - loss: 3.2252e-04  
Epoch 42/50  
30/30 [=====] - 2s 70ms/step - loss: 2.7601e-04  
Epoch 43/50  
30/30 [=====] - 2s 68ms/step - loss: 3.0736e-04  
Epoch 44/50  
30/30 [=====] - 2s 53ms/step - loss: 2.6299e-04  
Epoch 45/50  
30/30 [=====] - 2s 53ms/step - loss: 3.0173e-04  
Epoch 46/50  
30/30 [=====] - 2s 53ms/step - loss: 3.2376e-04  
Epoch 47/50  
30/30 [=====] - 2s 53ms/step - loss: 3.0133e-04  
Epoch 48/50  
30/30 [=====] - 2s 52ms/step - loss: 3.2601e-04  
Epoch 49/50  
30/30 [=====] - 2s 53ms/step - loss: 2.1707e-04  
Epoch 50/50

```
30/30 [=====] - 2s 52ms/step - loss: 3.1730e-04
```

```
<keras.callbacks.History at 0x7f54fd20dad0>
```

```
GRU_predicted_stock_price = model_gru.predict(X_test)
GRU_predicted_stock_price =
sc.inverse_transform(GRU_predicted_stock_price)
plot_predictions(test_set, GRU_predicted_stock_price)
```




```
return_rmse(test_set,GRU_predicted_stock_price)
```

The root mean squared error is 7.46.


## 5. Experiments Report

*Provide a summary of experimental results (e.g. tables, figures, charts).*


*Explain the meaning of your result and how your model can be used to address the related business problem.*



## LSTM MODEL



## GRU MODEL



The above results clearly show that the GRU model performed better than LSTM, with a similar structure and hyperparameters.

Above we tried to solve the business problem which is to predict the stock prices of the Master card using two different recurrent neural network models one is LSTM and Other one is GRU model.

The first graph shows the data plot of the actual dataset from 2016 to 2021. And later With the LSTM model we can see that the root mean square error is 9.24 which is a bit hgiht compared to the other model which is GRU model. GRU model got the root mean squared error of 7.46

References: <https://www.datacamp.com/tutorial>